




MIXED PRECISION TRAINING FOR IMAGE INPAINTING

Guilin Liu

IMAGE INPAINTING

Applications

Image Inpainting: remove unwanted content and re-fill with reasonable content.



MODEL DETAILS

- Partial Convolutions
- U-Net architecture with skip connections
- VGG loss (feature loss and gram matrix loss)

OUR CONTRIBUTION - PARTIAL CONVOLUTIONS



X : input

$$\begin{array}{ccc} 1 & 1 & 1 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \end{array}$$

M : mask

- Convolution (typical)

$$\text{Pixel update: } X'_i = W^T X + b$$

K : conv kernel size
normalization

- Partial Convolution

$$\text{Pixel update: } X'_i = W^T(X \circ M) \cdot \frac{K^2}{\text{sum}(M)} + b$$

$$\text{Mask update: } M'_i = \begin{cases} 1 & \text{if } \text{sum}(M) > 0 \\ 0 & \text{if } \text{sum}(M) = 0 \end{cases}$$



mask updating after several partial conv layers

As the receptive field becomes larger, mask will become all 1

OUR SOLUTION: PARTIAL CONVOLUTION

- Advantages
 - Initial value in holes doesn't matter
 - U-Net structure is possible
 - Another padding scheme
 - Possibly useful for other incomplete data

CHANGES REQUIRED

```
from apex import amp
```


```
self.generator, self.optimizer = amp.initialize(self.generator,  
                                              self.optimizer, opt_level=args.fp16_opt_level)
```

```
self.vgg_feat_loss = amp.initialize(self.vgg_feat_loss,  
                                     opt_level=args.fp16_opt_level)
```

```
with amp.scale_loss(total_loss, self.optimizer) as scaled_loss:  
    scaled_loss.backward()
```

GRAM MATRIX COMPUTATION

$$G_{ij}^l = \sum_k F_{ik}^l F_{jk}^l.$$



Source: <https://www.deeplearning.ai/>

EFFICIENT GRAM MATRIX COMPUTATION

Same code implementation for FP32/FP16

- More performant version for batch matrix-matrix product of matrices
 - **baddbmm** $\text{out}_i = \beta \text{ input}_i + \alpha (\text{batch1}_i @ \text{batch2}_i)$
 - saves memory traffic on reads/writes
- Factor `1./(c * h * w)` can be too small for FP16
 - large channel count or image resolution, e.g. when value less than 2^{-24}
 - change one-step big division to two-step smaller divisions to avoid underflow

```
def gram_matrix(input_tensor):
    """Compute gram matrix."""
    b, c, h, w = input_tensor.size()
    features = input_tensor.view(b, c, h * w)
    features_t = features.transpose(1, 2)
    inputs = torch.zeros(b, c, c).type(features.type())
    gram = torch.baddbmm(inputs, features, features_t,
                         beta=0, alpha=1./(c * h * w), out=None)
    return gram
```

DEALING WITH MACHINE EPSILON

- Meant to avoid division by zero
 - make your small constant number a bit bigger for FP16
 - e.g. **1e-8 → 1e-6**


```
1 self.mask_ratio = self.slide_winsize/(self.update_mask + 1e-8)
```

```
1 self.mask_ratio = self.slide_winsize/(self.update_mask + 1e-6)
```

OBTAINED SPEEDUPS


1.4x faster (~1.8x faster with doubled batch size)

Relative samples per second wrt w/o AMP




EFFECT ON ACCURACY

Same SSIM scores



METRICS OVER TIME



TAKEAWAYS

1. AMP is easy to integrate with very few changes
 - only 4 lines of code
 - use baddbmm for perf and watch out for too small normalization factor
 - avoid division by zero in reduced precision, e.g. $1e-8 \rightarrow 1e-6$
2. For the same model and batch size, we can get 1.4x speed up
3. For 2x larger batch size, we can get 1.8x speed up



NVIDIA®

